



targets: A package for reproducible workflows in R

Jesse Harrison
Analytics Safari (16.04.2021)



Common problem-solving scenario in R

```


1 # Structure of a typical linear R script
2
3 # 1. Load packages
4
5 library(x)
6
7 # 2. Load data from a CSV file
8
9 read_csv(...)
10
11 # 3. Clean + inspect the data
12 # (e.g. remove NAs)
13
14 cleandata <- mutate(...)
15 ggplot(...)
16
17 # 4. Do some modelling
18
19 model <- biglm(...)|

```

- 1) Launch a script
- 2) Wait while it runs
- 3) Find out there's an issue
- 4) Back to square one

Delays

A solution: pipeline toolkits



Take a **dependency graph** of the entire workflow

Skip steps (or **targets**) that are unchanged since the last run

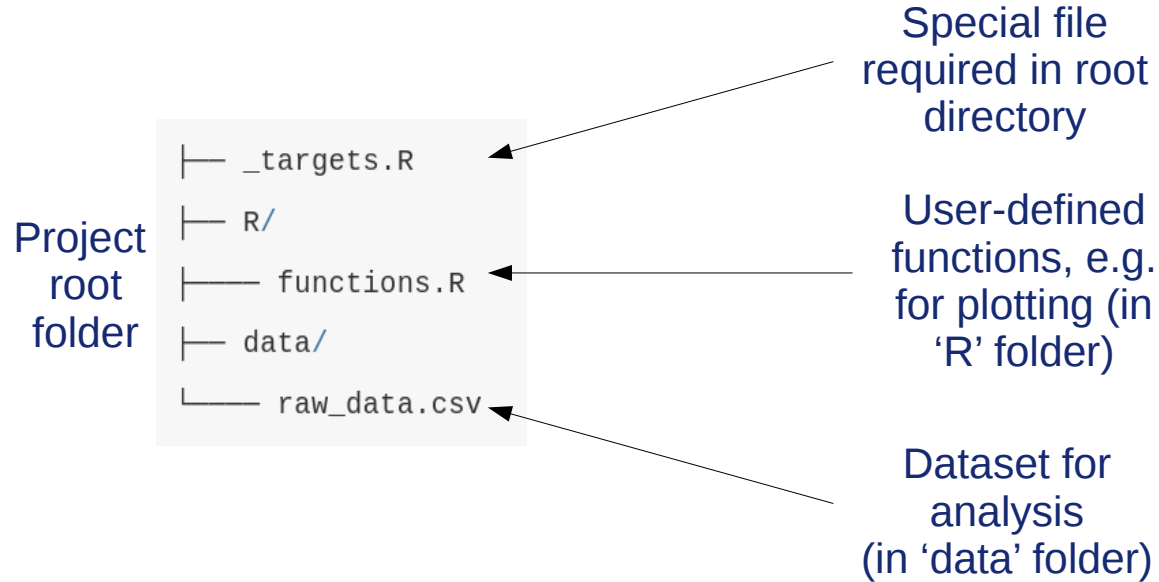
→ When all targets are up-to-date, can be confident that results match the underlying code and data

→ Additionally, can afford significant runtime speed-ups

targets: a pipeline toolkit package for R (and part of the ‘Targetopia’)



How it works



How it works

targets.R:

- Loads targets and other required packages
- Loads user-defined functions
- Sets default settings for all targets (e.g. data storage formats)
- Defines individual targets (i.e. intermediate workflow steps)

```
# _targets.R file
library(targets)
tar_option_set(packages = c("biglm", "dplyr", "ggplot2", "readr"))
```

Packages

```
# Most workflows have custom functions to support the targets.
read_clean ← function(path) {
  path %>%
    read_csv(col_types = cols()) %>%
    mutate(Ozone = replace_na(Ozone, mean(Ozone, na.rm = TRUE)))
}
```

```
fit_model ← function(data) {
  biglm(Ozone ~ Wind + Temp, data)
}
```

Functions

(can also be given directly in _targets.R)

```
create_plot ← function(data) {
  ggplot(data) +
    geom_histogram(aes(x = Ozone), bins = 12) +
    theme_gray(24)
}
```


```
# List of targets.
```

```
list(
  # airquality dataset in base R:
  tar_target(raw_data_file, "raw_data.csv", format = "file"),
  tar_target(data, read_clean(raw_data_file)),
  tar_target(fit, fit_model(data)),
  tar_target(hist, create_plot(data))
)
```

Targets (listed at the end)

`tar_visnetwork()`

For interactive dependency graphs



`tar_make()` runs the correct targets in the correct order.

R console

`tar_make()`

`#> ● run target raw_data_file`

`#> ● run target data`


`#> ● run target fit`

`#> ● run target hist`

`#> ● end pipeline`

<https://ropensci.org/blog/2021/02/03/targets/>

tar_visnetwork() For interactive dependency graphs



The next time you run `tar_make()`, `targets` skips everything that is already up to date, which saves a lot of time in large projects with long runtimes.

```
tar_make()  
#> ✓ skip target raw_data_file  
#> ✓ skip target raw_data  
#> ✓ skip target data  
#> ✓ skip target fit  
#> ✓ skip target hist  
#> ✓ skip pipeline
```

<https://ropensci.org/blog/2021/02/03/targets/>

Can specify no.s
of workers

Parallelism in cluster environments

Alternatives `tar_make_clustermq()` and `tar_make_future()` leverage `clustermq`⁶ and `future`⁷, respectively, to distribute targets on traditional schedulers such as `SLURM`⁸. It is only a matter of time before these backends become capable of `sending jobs to the cloud`⁹.

Storing results in S3 buckets (default is `_targets/`)

```
# Example _targets.R
library(targets)
tar_option_set(resources = list(bucket = "my-test-bucket-25edb4956460647d"))
write_mean <- function(data) {
  tmp <- tempfile()
  writeLines(as.character(mean(data)), tmp)
  tmp
}
list(
  tar_target(data, rnorm(5), format = "aws_qs"),
  tar_target(mean_file, write_mean(data), format = "aws_file")
)
```

Specify a bucket

Indicate AWS in target list

Machine learning with targets

<https://github.com/wlandau/targets-keras>

targets R package Keras model example

RStudio Cloud

The goal of this workflow is find the Keras model that best predicts customer attrition (“churn”) on a subset of the [IBM Watson Telco Customer Churn dataset](#). (See [this RStudio Blog post by Matt Dancho](#) for a thorough walkthrough of the use case.) Here fit multiple Keras models to the dataset with different tuning parameters, pick the one with the highest classification test accuracy, and produce a trained model for the best set of tuning parameters we find.

Newest Puhti r-env-singularity module (4.0.4)
comes with CUDA libraries + R interface to TensorFlow

<https://docs.csc.fi/apps/r-env-singularity/#r-interface-to-tensorflow>

Testing the targets Keras example on Puhti

```
#!/bin/bash
#SBATCH --account=dac
#SBATCH --output=output_%j.txt
#SBATCH --error=errors_%j.txt
#SBATCH --mail-type=END
#SBATCH --mail-user=jesse.harrison@csc.fi
#SBATCH --partition=test
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=8
#SBATCH --mem=30G
#SBATCH --time=00:15:00

ml r-env-singularity/4.0.4

# Clean up .Renviron file in home directory
if test -f ~/.Renviron; then
    sed -i '/TMPDIR/d' ~/.Renviron
    sed -i '/OMP_NUM_THREADS/d' ~/.Renviron
fi

# Specify a temp folder path
echo "TMPDIR=/scratch/dac/harrison" >> ~/.Renviron

# Match thread and core numbers
echo "OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK" >> ~/.Renviron

srun singularity_wrapper exec Rscript --no-save keras-targets.R
```

```
setwd('/scratch/dac/harrison/targets/keras')

library(targets)
tar_make()
```

```
v skip target units
v skip target act
v skip target data_file
v skip target data
v skip target recipe
* run branch run_e5d4783a
* run branch run_7f2173b5
* run branch run_c2634968
* run branch run_7c1ebb28
* run target best_run
* run target best_model
* run target report
* end pipeline
```

Needed to
install
some R
packages
and then
re-run

Target factories: making life easier for users

- Pre-packaged function sets that abstract away lower-level code
- For specific / specialised purposes

```
# _targets.R file
library(targets)
library(example.package)
biglm_factory("raw_data.csv")
```

What the user sees
(three lines!)

```
# script inside example.package

# Desport
read_csv <- function(path) {
  path %>%
    read_csv(col_types = cols()) %>%
    mutate(Ozone = replace_na(Ozone, mean(Ozone, na.rm = TRUE)))
}

# Desport
fit_model <- function(data) {
  biglm(Ozone ~ Wind + Temp, data)
}

# Desport
create_plot <- function(data) {
  ggplot(data) +
    geom_histogram(aes(x = Ozone, bins = 12) +
      theme_gray())
}

# Description (concise shorthand to express our example pipeline.
# Details
# Target factories should use "tar_target_raw()".
# "tar_target()" is for users, and "tar_target_raw()" is for developers.
# The former quotes its arguments, while the latter evaluates them.
# Desport
biglm_factory <- function(file) {
  list(
    tar_target_raw("raw_data_file", as.expression(file), format = "file"),
    tar_target_raw("data", quote(example.package::read_csv(raw_data_file))),
    tar_target_raw("fit", quote(example.package::fit_model(data))),
    tar_target_raw("hist", quote(example.package::create_plot(data)))
  )
}
```

All the magic the
user doesn't see

Situations where this could be useful

Further ideas / suggestions welcome!

- Jobs with lots of intermediate steps
- Analyses with bottlenecks
- I/O-intensive jobs

Less useful for jobs involving a single big analysis that runs for days...

... but for those, less need for a pipeline toolkit anyway!



repo status **Active** targets is the core engine of the targetopia. It learns the components of your data analysis project, runs the work with distributed computing, and skips steps that are already up to date. It reduces the runtime of successive runs, and it shows tangible evidence that your results match the underlying code and data.



repo status **Active** tarchetypes makes it easy to add certain kinds of common tasks to reproducible pipelines. Most of its functions create families of targets for parameterized R Markdown, simulation studies, and other general-purpose scenarios.



repo status **Active** stantargets is a workflow framework for Bayesian data analysis with cmdstanr. With concise, easy-to-use syntax, it defines versatile families of targets tailored to Bayesian statistics, from a single MCMC run with postprocessing to large simulation studies.



repo status **Active** Like stantargets, jagstargets is a workflow framework for Bayesian data analysis, with support for both single MCMC runs and large-scale simulation studies. It invokes JAGS through the R2jags package, which has nice features such as the ability to parallelize chains across local R processes.



brmstargets

repo status **Concept** brmstargets is an idea first proposed here. An implementation is planned, but no work has started. The goal is to accommodate brms-powered Bayesian data analysis workflows just as stantargets enhances cmdstanr.

Other ideas

Following precedent of stantargets, it should be possible to extend the R Targetopia to more methodology packages whose users face intense computation, long runtimes, and rapid changes. Possibilities include greta, nimble, keras, torch, torchvision, tidymodels, mlr3, and nlmixr.